



High Performance
Computing &
Big Data Services



hpc.uni.lu



hpc@uni.lu



[@ULHPC](https://twitter.com/ULHPC)

HPC School - Beginner

S1-2 - Don't fear the command line



Before we start



Objectives:

- get acquainted to the linux command line interface (CLI)
- be able to manipulate the file system
- be able to decrypt complicated commands

Prerequisite: you should be able to connect to the HPC cluster

A little bit of history

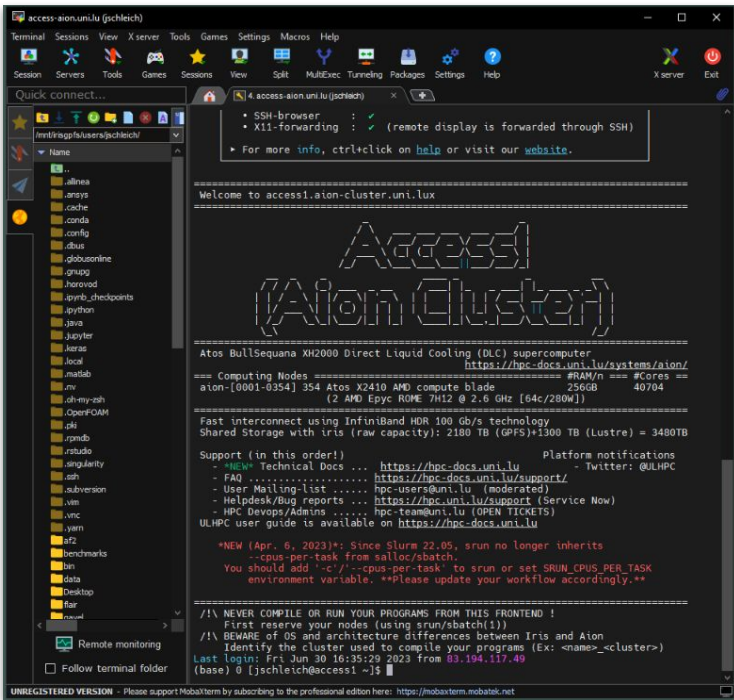
- Computers before the early 80s'
 - room sized, expensive
 - as powerful as a modern scientific calculator
 - multi user
 - central computer / terminal model
- Philosophy
 - use resources as efficiently as possible
 - textual human/machine interface (shell)
 - small specialized programs

Similarities with the modern HPC and Cloud ethos



Copyright © 2015 Pearson Education, Inc. All rights reserved.

Mobaxterm

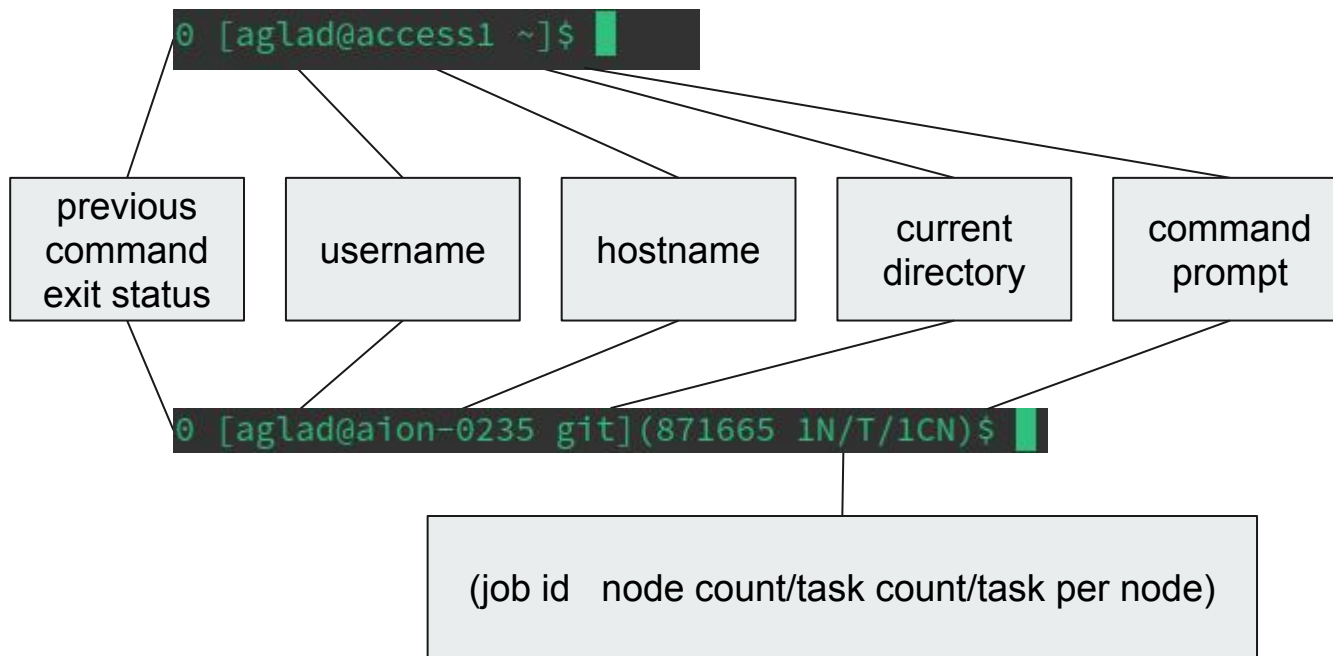


CLI



The command prompt

The bash shell



The linux file system

- Tree structure starting at /
 - No drives as in windows (c:, ...)
- Strong conventions
 - System directories
 - /etc - configuration files
 - /bin - built-in programs
 - /usr - 'user system resource'
other installed programs, libraries, ...
 - /home - users personal directories
 - Program installation follows conventions
 - difficult to know what is installed
 - the system able to provide **completion**
- Files and directories
 - ~ - shorthand for /home/users/<user_name>
 - . - current directory
 - .. - parent directory
 - .filename - hidden file/directory

```
total 512
drwx----- 18 hpcuser hpcuser 4096 Sep  7 15:25 .
drwxr-xr-x  5 root    root    4096 Jul 21 09:41 ..
-rw-----  1 hpcuser hpcuser   59 Jul 21 19:26 .bash_history
-rw-r--r--  1 hpcuser hpcuser   21 Jan  9 2022 .bash_logout
-rw-r--r--  1 hpcuser hpcuser   57 Jan  9 2022 .bash_profile
-rw-r--r--  1 hpcuser hpcuser 3824 Jan 14 2022 .bashrc
drwxr-xr-x 13 hpcuser hpcuser 4096 Aug  7 09:55 .cache
drwxr-xr-x 14 hpcuser hpcuser 4096 Aug  7 09:54 .config
-rw-r--r--  1 hpcuser hpcuser    0 Jul 21 11:12 data.dat
drwxr-xr-x  2 hpcuser hpcuser 4096 Aug  7 09:51 Desktop
-rw-r--r--  1 hpcuser hpcuser 4855 Oct 30 2017 .dir_colors
drwxr-xr-x  2 hpcuser hpcuser 4096 Aug  7 09:51 Documents
drwxr-xr-x  2 hpcuser hpcuser 4096 Aug  7 09:51 Downloads
```

- Access rights
 - user - group
 - d[rwx][rwx][rwx]
 - directory [user][group][other]
 - read - write - execute/navigate

Anatomy of a command



program [flags]... [arguments]...

- flags = options that change the behavior of the program
 - not sensitive to order
 - long flags - more readable, annoying to type
 - start with -- e.g. `--all`
 - can have parameters e.g. `--format=long --ignore foo`
 - short flags - condensed, difficult to parse
 - start with - e.g. `-l -a`
 - can be combined - e.g. `-la`. The order of the tags is not relevant `-la = -al`
 - can also have parameters - e.g. `-l foo`
- arguments = parameters of the program
 - positional
 - number of arguments depends on the program
- example - all these commands are equivalent
 - `$ls -la`
 - `$ls -a -l`
 - `$ls --format=long -a`

A little bit of help

- Manual page ← *more comprehensive*
 - `man <command>`
 - e.g. `$ man ls`
- Help flag ← *when a man page does not exist*
 - `<command> --help`
 - e.g. `$ ls --help`

```
LS(1)                                User Commands                                LS(1)

NAME
    ls - list directory contents

SYNOPSIS
    ls [OPTION]... [FILE]...

DESCRIPTION
    List information about the FILES (the current directory by default).
    Sort entries alphabetically if none of -cftuvSUX nor --sort is speci-
    fied.

    Mandatory arguments to long options are mandatory for short options
    too.

    -a, --all
        do not ignore entries starting with .

    -A, --almost-all
        do not list implied . and ..

    --author
        with -l, print the author of each file

    -b, --escape
        print C-style escapes for nongraphic characters

Manual page ls(1) line 1 (press h for help or q to quit)
```


A little bit of help

Now, you do it!

- What does the `echo` command do?
- What does the `which` command do?



A little bit of help

Now, you do it!

- What does the **echo** command do?
- What does the **which** command do?



ECHO(1)

NAME

echo - display a line of text

SYNOPSIS

```
echo [SHORT-OPTION]... [STRING]...  
echo LONG-OPTION
```

DESCRIPTION

Echo the STRING(s) to standard output.

WHICH(1)

General Commands Manual

WHICH(1)

NAME

which - shows the full path of (shell) commands.

SYNOPSIS

```
which [options] [--] programname [...]
```

DESCRIPTION

Which takes one or more arguments. For each of its arguments it prints to stdout the full path of the executables that would have been executed when this argument had been entered at the shell prompt. It does this by searching for an executable or script in the directories listed in the environment variable **PATH** using the same algorithm as **bash(1)**.

Lätz build ourselves a little playground



Go to the home directory

```
$ cd
```

Pull the repository containing the files

```
$ git clone https://gitlab.uni.lu/hlst/hpc-school-for-beginners.git
```

Check the content

You should see a couple of files and directories. Check that the CLI directory is present.

```
$ ls ~/hpc-school-for-beginners
```

Go to the directory containing the files for the tutorial

```
$ cd hpc-school-for-beginners/CLI
```

Navigating through the file system

Tab autocompletes paths and your commands



- **pwd** - print working directory
 - show the full path of the current directory
 - useful to know where you are
- **ls** - list
 - list the files and directories in the current directory
 - add a path in argument to show the content of another directory
 - **-a** flag shows hidden files and directories
 - **-l** formats the output and shows access right
- **cd** - change directory
 - **cd** with no argument returns you to your home dir
 - **cd /<path>** - go to the indicated absolute path
 - e.g. \$ **cd /home/users/hpcuser/foo/**
 - **cd ./<path>** or **cd <path>** - go to the relative path
 - e.g. \$ **cd ./foo** then \$ **cd nestedFoo**
 - **cd ..** - go to the parent directory
 - e.g. from ~/foo/nestedFoo -& **cd ../../dir**

```
[hpcuser@hpcschool git]$ pwd
/home/hpcuser/git
```

```
[hpcuser@hpcschool CLI]$ ls -la
total 24
drwxr-xr-x 6 root root 4096 Sep  8 14:27 .
drwxr-xr-x 4 root root 4096 Sep  8 14:27 ..
drwxr-xr-x 2 root root 4096 Sep  8 14:27 docs
```

```
[hpcuser@hpcschool CLI]$ cd docs
[hpcuser@hpcschool docs]$
```

Navigating through the file system

Now, you do it!

- go to your home directory
- from there, go to the tutorial directory
hpc-school-for-beginners/CLI/playground
- go back a level then to the docs directory



```
[hpcuser@hpcschool CLI]$ tree -d
.
├── docs
├── final_boss
├── playground
│   ├── backup
│   ├── files
│   │   └── experiments
│   │       └── data
│   ├── scripts
│   └── temp
```

Navigating through the file system

Now, you do it!



- go to your home directory
- from there, go to the tutorial directory
hpc-school-for-beginners/CLI/playground
- go back a level then to the docs directory

```
[hpcuser@hpcschool CLI]$ tree -d
├── docs
├── final_boss
├── playground
│   ├── backup
│   ├── files
│   │   └── experiments
│   │       └── data
│   ├── scripts
│   └── temp
```

```
[hpcuser@hpcschool git]$ cd
```

```
[hpcuser@hpcschool ~]$
```

```
[hpcuser@hpcschool ~]$ cd hpc-school-for-beginners/CLI/
playground/
```

```
[hpcuser@hpcschool playground]$
```

```
[hpcuser@hpcschool playground]$ cd ../docs/
```

```
[hpcuser@hpcschool docs]$
```

Executing programs and scripts



- Built-in and installed software
 - Built-in commands come with the shell
 - Installed software
 - via package manager or install scripts
 - copies files according to conventions
 - The OS is aware of their existence/location
 - just type their name (e.g. `ls`, `cd`)
 - the tab key proposes completion (e.g. `l + tab` -> all executables starting with 'l')
- Scripts and non installed software
 - Launcher scripts (`.sh`), precompiled software from archives
 - The OS is usually **not** aware of their existence
 - must be executable
 - called using its *path/name* (either absolute (starts with `/`) or relative (starts with `./`))

Executing programs and scripts

Now, you do it!



- Go to `~/hpc-school-for-beginners/CLI`
- Run the backup script located in `playground/scripts/backup.sh` with the file `playground/files/important.txt` as an argument
- Check the content of the `playground/backup` directory

backup.sh is a custom script and does not have a man page or --help flag.
Usage is `$./backup.sh <file_to_backup>`



Executing programs and scripts

Now, you do it!



- Go to `~/hpc-school-for-beginners/CLI`
- Run the backup script located in `playground/scripts/backup.sh` with the file `playground/files/important.txt` as an argument
- Check the content of the `playground/backup` directory

backup.sh is a custom script and does not have a man page or --help flag.
Usage is `$./backup.sh <file_to_backup>`



```
[hpcuser@hpcschool ~]$ cd hpc-school-for-beginners/CLI
[hpcuser@hpcschool CLI]$ ls -l playground/backup/
total 0
[hpcuser@hpcschool CLI]$ ./playground/scripts/backup.sh
playground/files/important.txt
[hpcuser@hpcschool CLI]$ ls -l playground/backup/
total 4
-rw-r--r-- 1 hpcuser hpcuser 20 Sep 11 16:32 important.
txt-230911-16:32.bak
[hpcuser@hpcschool CLI]$
```

Manipulating files (1/2)

- **mkdir - make directory**

- create a directory at the designated path
 - `$ mkdir test`
 - `$ mkdir i_dont_exist/test` -> The command fails because the `i_dont_exist` directory does not exist.
 - `$ mkdir -p i_dont_exist/test` -> Recursively creates the directories if they do not exist.



- **cp - copy**

- copy a file - `cp <source> <destination>`
 - `$ cp ./dir/file.txt file(copy).txt`
 - `$ ls` -> you should see test.txt
- copy a directory - use the `-r` flag
 - `$ cp -r ./dir ./foo/nestedFoo`
 - `$ cp -r ./dir ./foobar`
The command fails. The foobar directory doesn't exist.
- copy files using a pattern
 - `$ cp dir/* foo`
copies all files in dir to foo
 - `$ cp dir/*.txt foo`
copies all files ending in .txt in dir to foo

Manipulating files (1/2)

Now, you do it!



- Create a 'manual_backup' directory in the CLI directory
- Make a backup of
~/hpc-school-for-beginner/CLI
/playground/temp/experiment.out in the
'manual_backup' directory

Manipulating files (1/2)

Now, you do it!



- Create a 'manual_backup' directory in the CLI directory
- Make a backup of
~/hpc-school-for-beginner/CLI
/playground/temp/experiment.out in the
'manual_backup' directory

```
[hpcuser@hpcschool CLI]$ mkdir manual_backup
[hpcuser@hpcschool CLI]$ ls
docs  final_boss  manual_backup  playground
[hpcuser@hpcschool CLI]$ cp playground/temp/experiment.
out ./manual_backup/experiment_backup.out
[hpcuser@hpcschool CLI]$ ls manual_backup/
experiment_backup.out
```

Manipulating files (2/2)

- **rm - remove file**

- There is no bin. Deleted files **cannot** be recovered.
- delete a file
 - `$ rm ~/hpcschool/foo/data.dat`
- delete a directory
 - `$ rm -r ~/hpcschool/foo/nestedFoo`
- force deletion
 - `$ rm -f -r ~/hpcschool/foo`

- **mv - move file**

- move a file/directory (cut and paste)
 - `$ mv test.txt foo/test.txt`
- rename a file/directory = moving a file to the same directory
 - `$ mv foo/test.txt foo/temp.txt`
- move and rename a file
 - `$ mv foo/temp.txt ../test.txt`



up and down arrows
navigate through recent
commands



Manipulating files (2/2)

Now, you do it!



- Delete the `~/hpc-school-for-beginner/CLI/playground/temp` directory and its content
- Move the backup of `experiment.out` to the `hpc-school-for-beginners/CLI/playground/files/experiment/data` directory and rename it `test001.com`

Manipulating files (2/2)

Now, you do it!



- Delete the ~/hpc-school-for-beginner/CLI/playground/temp directory and its content
- Move the backup of experiment.out to the hpc-school-for-beginners/CLI/playground/files/experiment/data directory and rename it test001.com

```
[hpcuser@hpcschool CLI]$ ls playground/
backup  files  scripts  temp
[hpcuser@hpcschool CLI]$ rm -r playground/temp
[hpcuser@hpcschool CLI]$ ls playground/
backup  files  scripts
[hpcuser@hpcschool CLI]$
```

```
[hpcuser@hpcschool CLI]$ ls playground/files/experiments/data
test003.com  test034.com  test072.com  test156.com
test013.com  test057.com  test077.com
test014.com  test061.com  test101.com
test021.com  test065.com  test121.com
[hpcuser@hpcschool CLI]$ ls manual_backup/
experiment_backup.out
[hpcuser@hpcschool CLI]$ mv manual_backup/experiment_backup.out playground/files/experiments/data/test001.com
[hpcuser@hpcschool CLI]$ ls playground/files/experiments/data
test001.com  test021.com  test065.com  test121.com
test003.com  test034.com  test072.com  test156.com
test013.com  test057.com  test077.com
test014.com  test061.com  test101.com
[hpcuser@hpcschool CLI]$ ls manual_backup/
[hpcuser@hpcschool CLI]$
```

Reading and writing files (1/2)

- **cat** - concatenates files and write to the standard output
 - `$ cat <filename>`
 - e.g. `$ cat dir/data.csv`
- **less** - reading longer files
 - `$ less <filename>`
 - e.g. `$ less dir/data.csv`
 - scroll with arrows/page up/page down
 - quit with q
- **tail** - show the last lines of a text file
 - `$ tail <filename>`
 - `$ tail -n 25 <filename>` - specify the number of displayed lines
 - `$ tail -f <filename>` - follow new lines
- **chaining commands** - the `|` operator
 - allow to pass the output of a command to the next one
 - `$ ls -la /usr/bin | less`
- **grep** - filtering utility
 - `$ ls /usr/bin | grep update`
 - `$ grep pattern filename`

ctrl-r allows to search your
command history



Reading and writing files (1/2)

Now, you do it!



- Launch `playground/scripts/tailMe.sh &` and follow the output file (`tailMe.out`)
- In `playground/files/experiments/data`, find all files that contain data about Methylene (with a capitalized M)

The `tailMe.sh` script writes data to the `tailMe.out` file every second for a minute.



Adding an `&` at the end of a command will make it run in the background. You will not see any output but it will not block the terminal while it is running.

Hit `ctrl + c` to interrupt any running program

Reading and writing files (1/2)

Now, you do it!



- Launch `playground/scripts/tailMe.sh` & and follow the output file (`tailMe.out`)
- In `playground/files/experiments/data`, find all files that contain data about Methylene (with a capitalized M)

The `tailMe.sh` script writes data to the `tailMe.out` file every second for a minute.



Adding an `&` at the end of a command will make it run in the background. You will not see any output but it will not block the terminal while it is running.

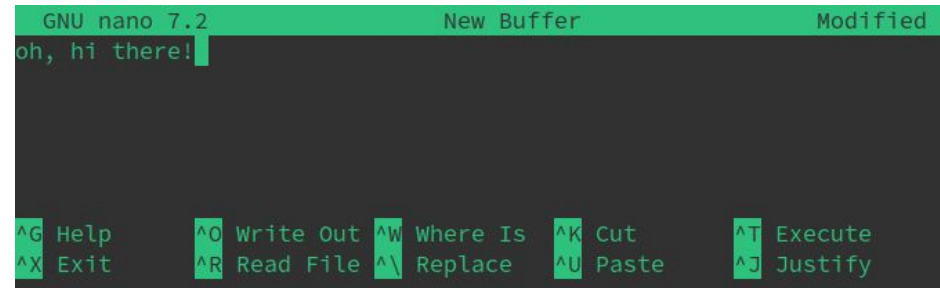
Hit `ctrl + c` to interrupt any running program

```
[hpcuser@hpcschool CLI]$ ./playground/scripts/tailMe.sh &
[1] 13509
[hpcuser@hpcschool CLI]$ ls
docs  final_boss  manual_backup  playground  tailMe.out
[hpcuser@hpcschool CLI]$ tail -f tailMe.out
processing 9/60
processing 10/60
processing 11/60
processing 12/60

[hpcuser@hpcschool data]$ cat test001.com | grep Methylene
Methylene uhf forces
[hpcuser@hpcschool data]$ grep Methylene test*.com
test001.com:Methylene uhf forces
test034.com:Methylene cisd/3-21g
test065.com:Methylene mp2 density, fully in-core
test065.com:Methylene mp2 density
test065.com:Methylene mp2 density, quartic out of core
test065.com:Methylene mp2 density, quintic out of core
test065.com:Methylene mp2 using l802/l901
test065.com:Methylene mp2 use=l903
test077.com:Methylene ump2 scan
```

Reading and writing files (2/2)

- > and >> - redirect the standard output (i.e. output of a command in the terminal) to a file
 - `& ls /usr/bin > all_bins.txt`
 - creates or opens in **overwrite** mode the `all_bins.txt` file and write the output of `ls /usr/bin` in it
 - `& ls -la /usr/bin >> all_bins.txt`
 - creates or opens in **append** mode the `all_bins.txt` file and write the output of `ls -la /usr/bin` in it
- the nano text editor
 - `$ nano [filename]`
 - basic usage
 - write with your keyboard
 - move the cursor with the arrow keys
 - shortcuts
 - at the bottom of the screen
 - accessible via `ctrl + <key>`
 - e.g. save: `ctrl + o`; quit: `ctrl + x`



```

GNU nano 7.2          New Buffer          Modified
oh, hi there!
^G Help    ^O Write Out ^W Where Is ^K Cut    ^T Execute
^X Exit    ^R Read File ^\ Replace  ^U Paste  ^J Justify
  
```

Changing permissions

chmod - **change file mode** bits
allows you to change access rights on
your files and directories

Why?

- give/restrict access to other people
- make some files executable

```
drwxr-xr-x  4 root    root    4096 Sep  8 14:27 .
drwx----- 19 hpcuser hpcuser 4096 Sep  8 14:30 ..
drwxr-xr-x  6 root    root    4096 Sep  8 14:27 CLI
drwxr-xr-x  8 root    root    4096 Sep  8 14:27 .git
-rw-r--r--  1 root    root    6213 Sep  8 14:27 README.md
```

\$ **chmod 744 <filename>**

- first number for the owner, second for the owner's group, third for everyone else
- octal mode
 - 1 = execute
 - 2 = write
 - 4 = read
- for directories
 - execute allows to cd into
 - write allows to create/delete files
 - read allows to see the content of the directory

sum the numbers to adjust the rights

- $7 = 1+2+4 = x + w + r$
- 4 = read only

Changing permissions

Now, you do it!

- make CLI/playground/files/secret.txt readable only to you
- make CLI/scripts/helloWorld.sh executable and run it



Changing permissions

Now, you do it!



- make CLI/playground/files/secret.txt readable only to you
- make CLI/scripts/helloWorld.sh executable and run it

```
[hpcuser@hpcschool files]$ ls -la secret.txt
-rw-r--r-- 1 hpcuser hpcuser 10 Sep 11 16:03 secret.txt
[hpcuser@hpcschool files]$ chmod 700 secret.txt
[hpcuser@hpcschool files]$ ls -la secret.txt
-rwx----- 1 hpcuser hpcuser 10 Sep 11 16:03 secret.txt
[hpcuser@hpcschool files]$
```

```
[hpcuser@hpcschool scripts]$ ls -l helloWorld.sh
-rw-r--r-- 1 hpcuser hpcuser 32 Sep 11 16:03 helloWorld.sh
[hpcuser@hpcschool scripts]$ ./helloWorld.sh
bash: ./helloWorld.sh: Permission denied
[hpcuser@hpcschool scripts]$ chmod 744 helloWorld.sh
[hpcuser@hpcschool scripts]$ ls -l helloWorld.sh
-rwxr--r-- 1 hpcuser hpcuser 32 Sep 11 16:03 helloWorld.sh
[hpcuser@hpcschool scripts]$ ./helloWorld.sh
hello world!
```

Moving data

- Rsync is a utility that allows to synchronize data between machines

- upload/download files
- synchronize files between servers
- resume interrupted transfers

- Push data

- `$ rsync -azvu <source> [user@]<host>:<destination>`
- `$ rsync -azvu data_directory aion-cluster:my_data`

- Pull data

- `$ rsync -azvu [user@]<host>:<source> <destination>`
- `$ rsync -azvu aion-cluster:my_data data_directory`

- Flags

- a - archive mode (recursive, copies files, rights, links, ...)
- z - compress data during transfer (speeds up transmission)
- v - verbose (display what is going on)
- u - update (skip files that are newer on the receiver)
- P - progress bar (monitor big transfers)

Now, you do it!



- copy the content of the CLI/docs directory to your machine to get this presentation and a command line interface cheat sheet pdf

Moving data

Now, you do it!

- copy the content of the CLI/docs directory to your machine to get this presentation and a command line interface cheat sheet pdf



Run rsync from your laptop, not from the HPC.
Finding the HPC cluster from your laptop is easier than finding your laptop from the cluster.



We need to add the **-e 'ssh -p 8022'** flag to access the custom ssh port of the cluster.

Moving data

Now, you do it!



- copy the content of the CLI/docs directory to your machine to get this presentation and a command line interface cheat sheet pdf

Run rsync from your laptop, not from the HPC.
Finding the HPC cluster from your laptop is easier than finding your laptop from the cluster.



We need to add the **-e 'ssh -p 8022'** flag to access the custom ssh port of the cluster.

```
aglad@hpcschool ~/tmp$ rsync -azvu -e 'ssh -p 8022' aglad@access-aion.uni.lu:~/hpc-school-for-b
eginners/CLI/docs/CLI_Cheat_Sheet.pdf ~/tmp
receiving incremental file list
CLI_Cheat_Sheet.pdf

sent 43 bytes  received 134,552 bytes  89,730.00 bytes/sec
total size is 175,255  speedup is 1.30
aglad@hpcschool ~/tmp$ ls ~/tmp
CLI_Cheat_Sheet.pdf
aglad@hpcschool ~/tmp$
```

Do the same for the other file or copy the whole directory at once.

Final Boss



```
$ man bash > tmp.dat
```

```
$ cat tmp.dat | grep -i bash | wc -l
```

Find out what these commands are doing. Don't run them yet!

Final Boss



```
$ man bash > tmp.dat
```

```
$ cat tmp.dat | grep -i bash | wc -l
```

Find out what these commands are doing. Don't run them yet!

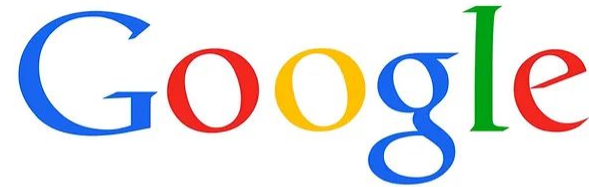
- first command
 - redirect the content of the man page command for the bash program to the tmp.dat file
- second command
 - display the content of the tmp.dat file (cat) and
 - pipe the result to grep. Only keep the lines that contain 'bash' while ignoring the case and
 - pipe the output to wc which will count the number of lines (-l flag)
- summary: count the number of lines containing 'bash' (case insensitive) in the man page of the bash program

Final Boss - Phase 2



```
$ export HPL_VERSION=2.3  
$ wget --continue http://www.netlib.org/benchmark/hpl/hpl-\${HPL\_VERSION}.tar.gz  
$ tar xvfz hpl-${HPL_VERSION}.tar.gz
```

Find out what these commands are doing? Don't run them yet!
You might need the help of google on this one!



Is your friend

Final Boss - Phase 2



```
$ export HPL_VERSION=2.3
$ wget --continue http://www.netlib.org/benchmark/hpl/hpl-2.3.tar.gz
$ tar xvfz hpl-2.3.tar.gz
```

Find out what these commands are doing? Don't run them yet!
You might need the help of google on this one!

- set an environment variable
- download a file from netlib.org. The name of the file depends on the value of the environment variable that has been set previously. If the file was partially downloaded, continue instead of redownloading everything.
- extract the files (x) from a gzip archive (z) in the hpl-2.3.tar.gz file (f) and show the logs (v)

Final Boss - Final form



Make `~/hpcschool1/data/runme.sh` executable and run it.

What did it do? How can you get rid of it? The script might contain clues and you have all the keys...

Final Boss - Final form



Make `~/hpcschool1/data/runme.sh` executable and run it.

What did it do? How can you get rid of it? The script might contain clues and you have all the keys...

- Too many files to be deleted one by one. Maybe they all have a pattern in common?
- It seems that you don't have the rights to remove files in this directory

Final Boss - Final form



Make `~/hpcschool/data/runme.sh` executable and run it.

What did it do? How can you get rid of it? The script might contain clues and you have all the keys...

- Too many files to be deleted one by one. Maybe they all have a pattern in common?
- It seems that you don't have the rights to remove files in this directory
- Regain write rights on the directory `$ chmod 700 ~/hpc-school-for-beginners/CLI`
- All files finish with a 1. Delete them using a pattern. `$rm -f ~/hpc-school-for-beginners/CLI/*1`

Never trust random scripts and commands found on the internet. Try to understand them first!

Your rights are limited and you cannot really hurt the HPC cluster.

You could easily lose you data however.